

Intro. to Computer Science :


Network System

May 2020

Honguk Woo

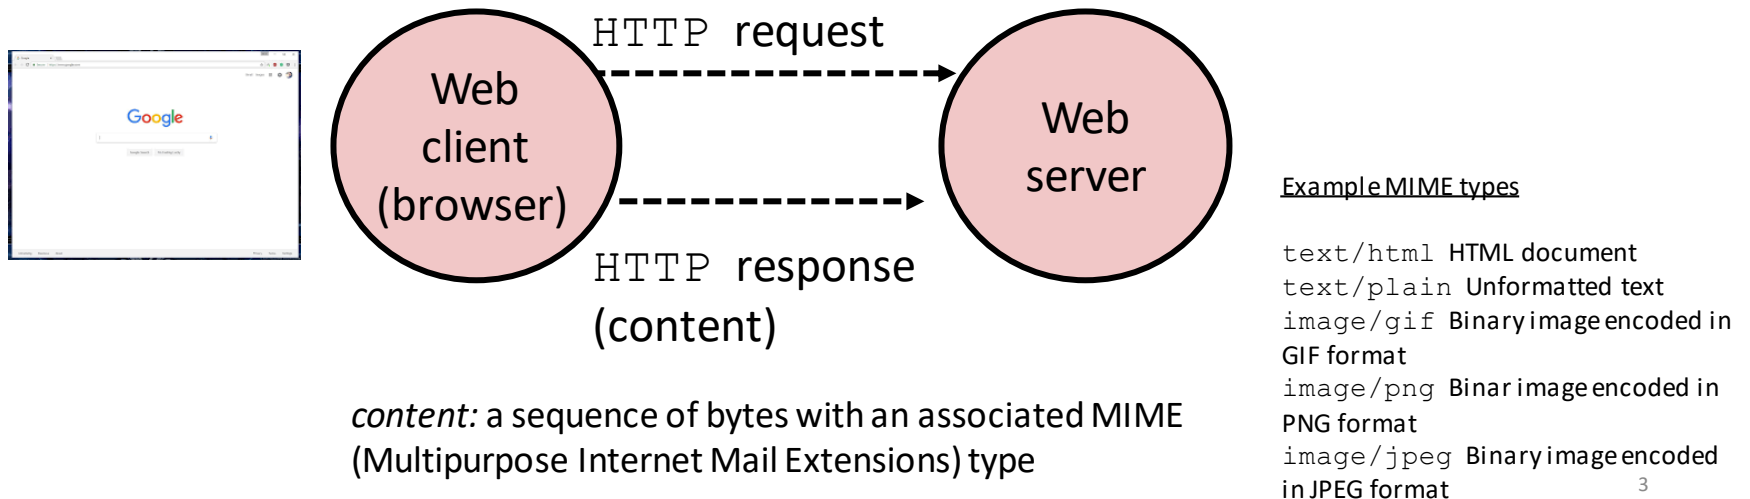
Client-Server Model

- Every network application is based on the client-server model:
 - A *server* process and one or more *client* processes
 - Server manages some *resource*
 - Server provides *service* by manipulating resource for clients
 - Clients and servers are processes running on hosts (can be the same or different hosts).



e.g., Web : Browser and Server

- Network Client and Server Model in Web Architecture
 - Web Browser : Google Chrome, MS explorer, Mozilla Firefox
 - Web Server : Web hosts for web services
 - Web server software : Apache HTTP server, MS IS, NGINX
 - Clients and servers communicate using the **HyperText Transfer Protocol (HTTP)**
 - Client and server establish TCP connection
 - Client requests content
 - Server responds with requested content
 - Client and server close connection (eventually)



e.g., Web : Static and Dynamic Content

- The content returned in HTTP responses can be either *static* or *dynamic*
 - *Static content*: content stored in files and retrieved in response to an HTTP request
 - Examples: HTML files, images, audio clips
 - Request identifies which content file
 - *Dynamic content*: content produced on-the-fly in response to an HTTP request
 - Example: content produced by a program executed by the server on behalf of the client
 - Request identifies file containing executable code
- *Web content is associated with a file that is managed by the server*

e.g., Web : Resources

- Unique name for a *file*: URL (Universal Resource Locator)
- Example URL: `http://www.host.edu:80/index.html`
 - Protocol: The application-level protocol used by the client and server, e.g., HTTP
 - Hostname: The DNS domain name of the server
 - Port: The TCP port number that the server is listening for incoming requests from the clients
 - Path-and-file-name: The name and location of the requested resource, under the server document base directory
- Use *prefix* (`http://www.host.edu:80`) to infer:
 - What kind (protocol) of server to contact (HTTP)
 - Where the server is (`www.host.edu`)
 - What port it is listening on (80)
- Use *suffix* (`/index.html`) to:
 - Find file on file system
 - Initial `"/` in suffix denotes home directory for requested content
 - Minimal suffix is `"/`, which server expands to configured default filename (usually, `index.html`)

e.g., Web : Resources – skku site

Chrome -> dev tool -> network

The screenshot shows the SKKU (Sungkyunkwan University) website in a Chrome browser. The page title is "성균관대학교 S-Gallery". The URL in the address bar is "https://www.skku.edu/skku/". The Chrome DevTools Network tab is open, displaying a list of resources. The resource "logo.png" is selected, and its details are shown in the right-hand pane. The details pane includes the following information:

- General:**
 - Request URL: `https://www.skku.edu/_res/skku//img/common/logo.png`
 - Request Method: GET
 - Status Code: 200 OK (from memory cache)
 - Remote Address: 1...
 - Referrer Policy: no-referrer-when-downgrade
- Response Headers:**
 - Accept-Ranges: bytes
 - Connection: Keep-Alive
 - Content-Length: 6790
 - Date: Mon, 20 May 2019 15:58:04 GMT
 - Keep-Alive: timeout=5, max=98
 - Last-Modified: Sat, 28 Jul 2018 13:06:46 GMT
- Request Headers:**
 - Provisional headers are shown
 - DNT: 1
 - Referer: `https://www.skku.edu/skku/`
 - User-Agent: Mozilla/5.0 (Windows NT 10.0; Win64; x64) AppleWebKit/537.36 (KHTML like Gecko) Chrome/74.0.3729.157 Safari/537.36

The bottom of the DevTools window shows a summary: "114 requests | 116 KB transferred | 8.2 MB resources | Finish: 12.68 s | D..."

e.g., Web : Resources – skku site

- An HTTP client sends an HTTP request to a server in the form of a request message which includes following format
 - Request-line : **GET** <https://www.skku.edu/skku/index.do> HTTP/1..1
 - Zero or more header fields
 - An empty line indicating the end of the header fields
 - Optionally a message-body
- *Try : How to view HTTP headers in Google Chrome?*
 - ① In Chrome, visit a URL, right click, select Inspect to open the developer tools
 - ② Select Network tab
 - ③ Reload the page, select any HTTP request on the left panel, and the HTTP headers will be displayed on the right panel
- <https://www.mkyong.com/computer-tips/how-to-view-http-headers-in-google-chrome/>

e.g., Web : sample HTTP request, response

- HTTP Request

GET /path/file.html HTTP/1.1

Host: www.host1.com:80

User-Agent: Mozilla/3.0

[blank line here]

- HTTP Response

HTTP/1.1 200 OK

Date: Fri, 31 Dec 201x 23:59:59 GMT

Content-Type: text/html

Content-Length: 1354

<html>

<body>

<h1>Happy New Millennium!</h1>

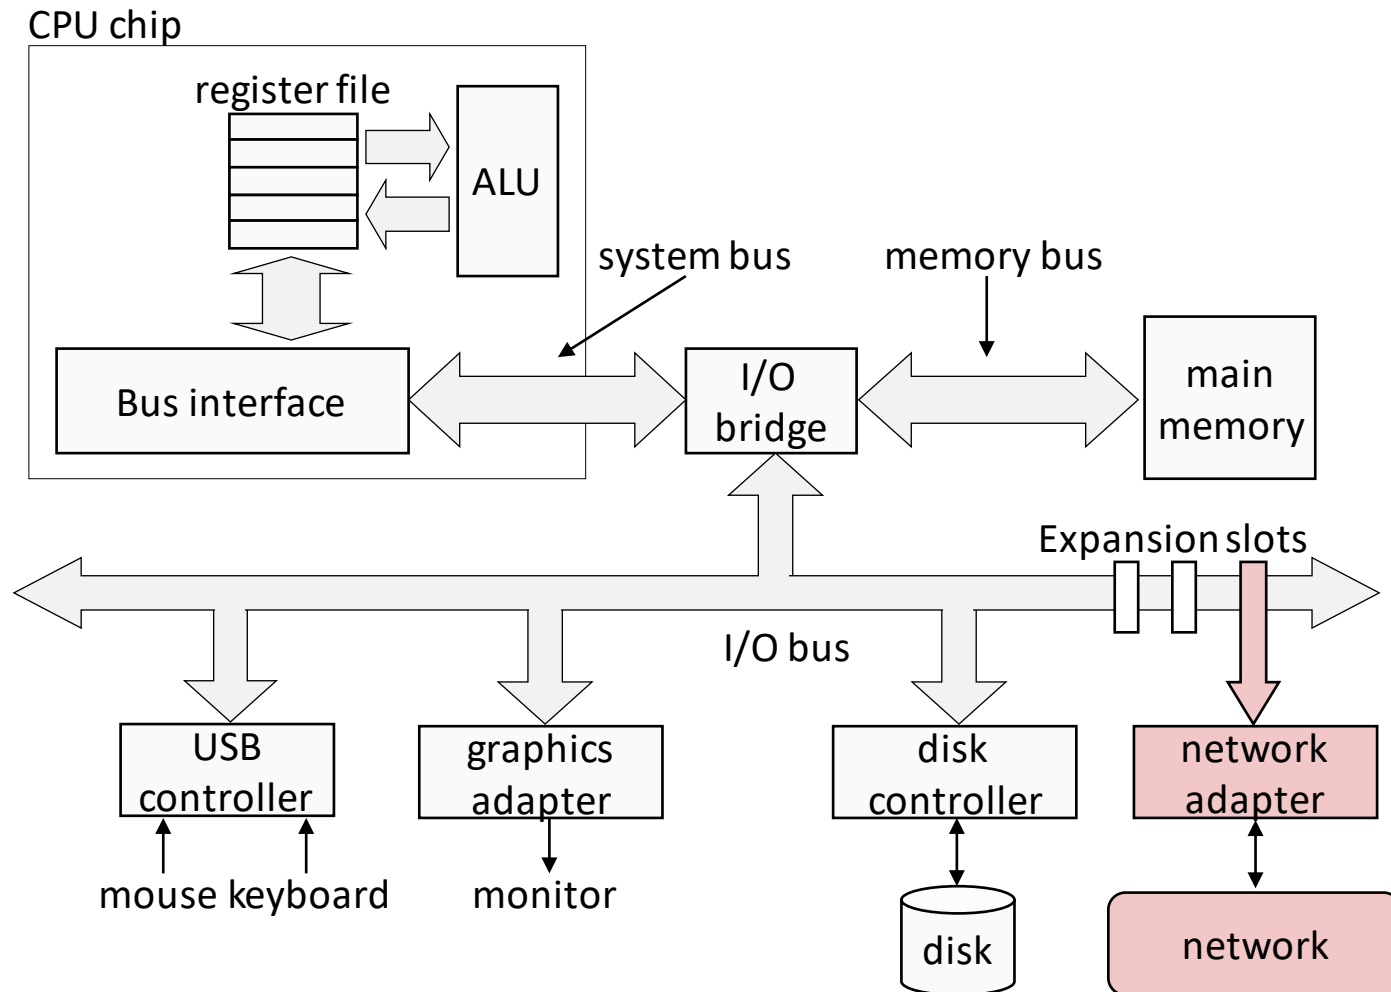
(more file contents) . . . </body>

</html>

History of Web


- <https://home.cern/science/computing/birth-web/short-history-web>
- <https://www.lopezferrando.com/a-brief-history-of-the-web/>

Hardware Organization of a Network Host



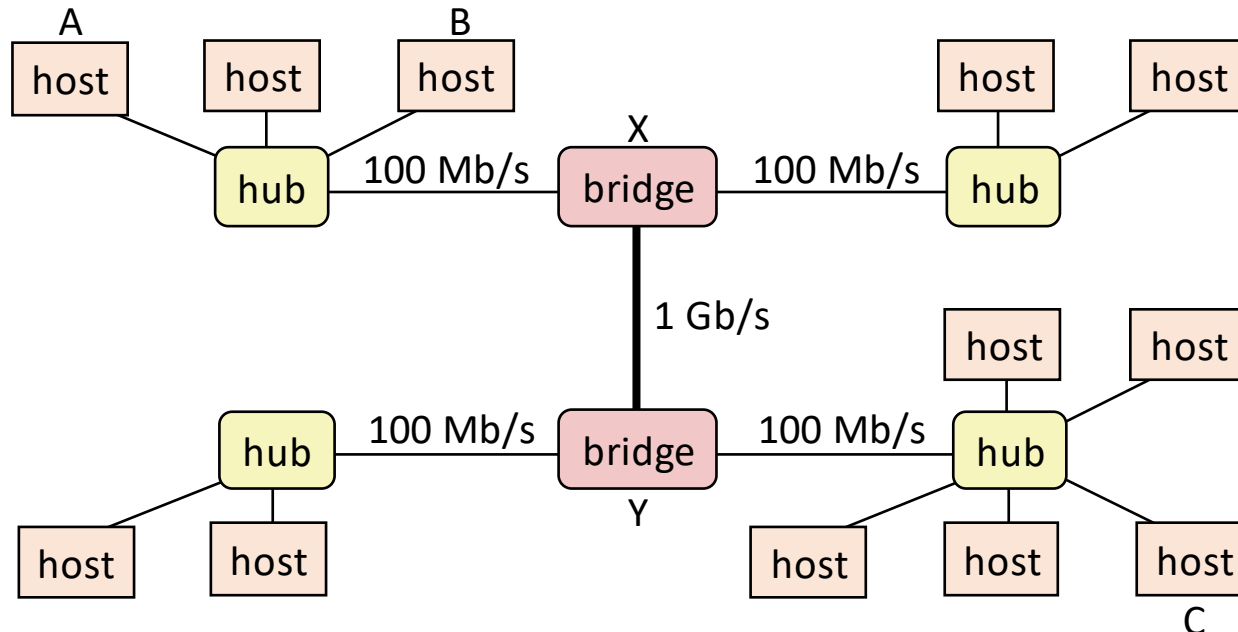
LAN (Local Area Network) - Ethernet Segment

- **Ethernet segment** consists of a collection of *hosts* connected by wires to a *hub*
- Spans room or floor in a building
- Operation
 - Each Ethernet adapter has a unique 48-bit address (MAC address)
 - E.g., 00:16:ea:e3:54:e6
 - Hosts send bits to any other host in chunks called *frames*
 - **Hub copies each bit from each port to every other port**
 - Every host sees every bit




LAN - Bridged Ethernet Segment

- Spans building or campus
- *Bridges* cleverly learn which hosts are reachable from which ports and then selectively copy frames from **port to port**
 - Bridge looks at the destination before forwarding unlike a hub




LAN – Conceptual View

- For simplicity, hubs, bridges, and wires are often shown as a collection of hosts attached to a single wire:




Next Level: internets

- Multiple incompatible LANs can be physically connected by specialized computers called *routers*
- The connected networks are called an *internet* (lower case)
 - *Internetwork(internet)*
 - The Global IP Internet (uppercase “I”) is the most famous example of an internet (lowercase “i”)
- *What’s AP (access point) at your home ?*




Logical Structure of an internet

- Ad hoc interconnection of networks
 - No particular topology
 - Vastly different router & link capacities
- Send packets from source to destination by hopping through networks
 - Router forms bridge from one network to another
 - Different packets may take different routes



Logical Structure of an internet : why ?


- Packet switching
 - Messages are broken into fixed-sized packets, Packets are routed independently, and the original message is re-assembled at the destination



Logical Structure of an internet : why ?


- A **decentralized** network with multiple paths between points A and B
 - Decentralized networks with redundant paths provide **robustness** in network design
 - Dividing the message into small packets that are routed independently
 - Each router along the path forwards packets to another router along the path

Centralized Network



Vulnerability

Decentralized Network



Robustness

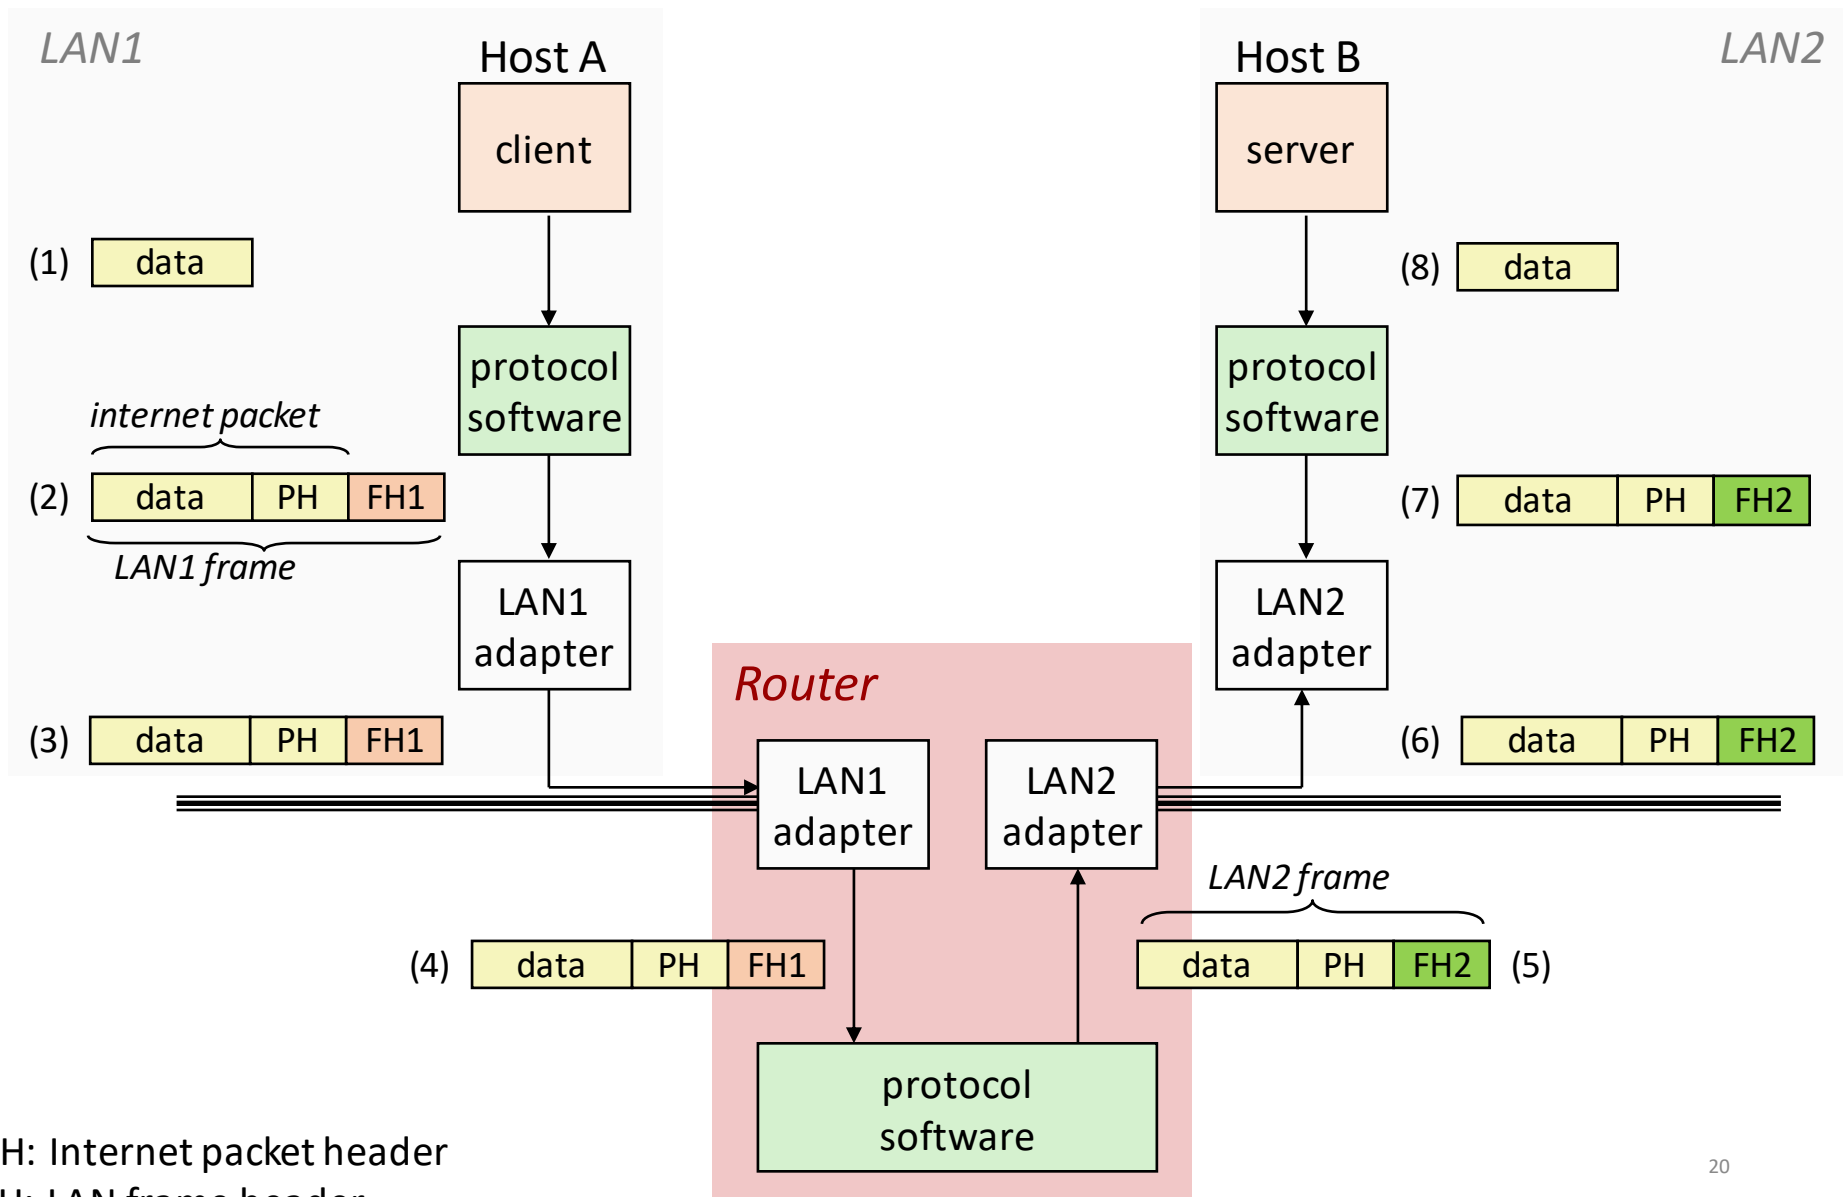
The Notion of an internet Protocol

- How is it possible to send bits across incompatible LANs and WANs?
- Solution: *protocol* software running on each host and router
 - Protocol is a set of rules that governs how hosts and routers should cooperate when they transfer data from network to network.
 - Smooths out the differences between the different networks

Role of Internet Protocol

- Provides a **naming** scheme
 - An internet protocol defines a uniform format for **host addresses**
 - Each host (and router) is assigned at least one of these internet addresses that uniquely identifies it
- Provides a **delivery** mechanism
 - An internet protocol defines a standard transfer unit (**packet**)
 - Packet consists of **header** and **payload**
 - Header: contains information such as packet size, source and destination addresses.
 - Payload: contains data bits sent from source host.

Transferring internet Data Via Encapsulation

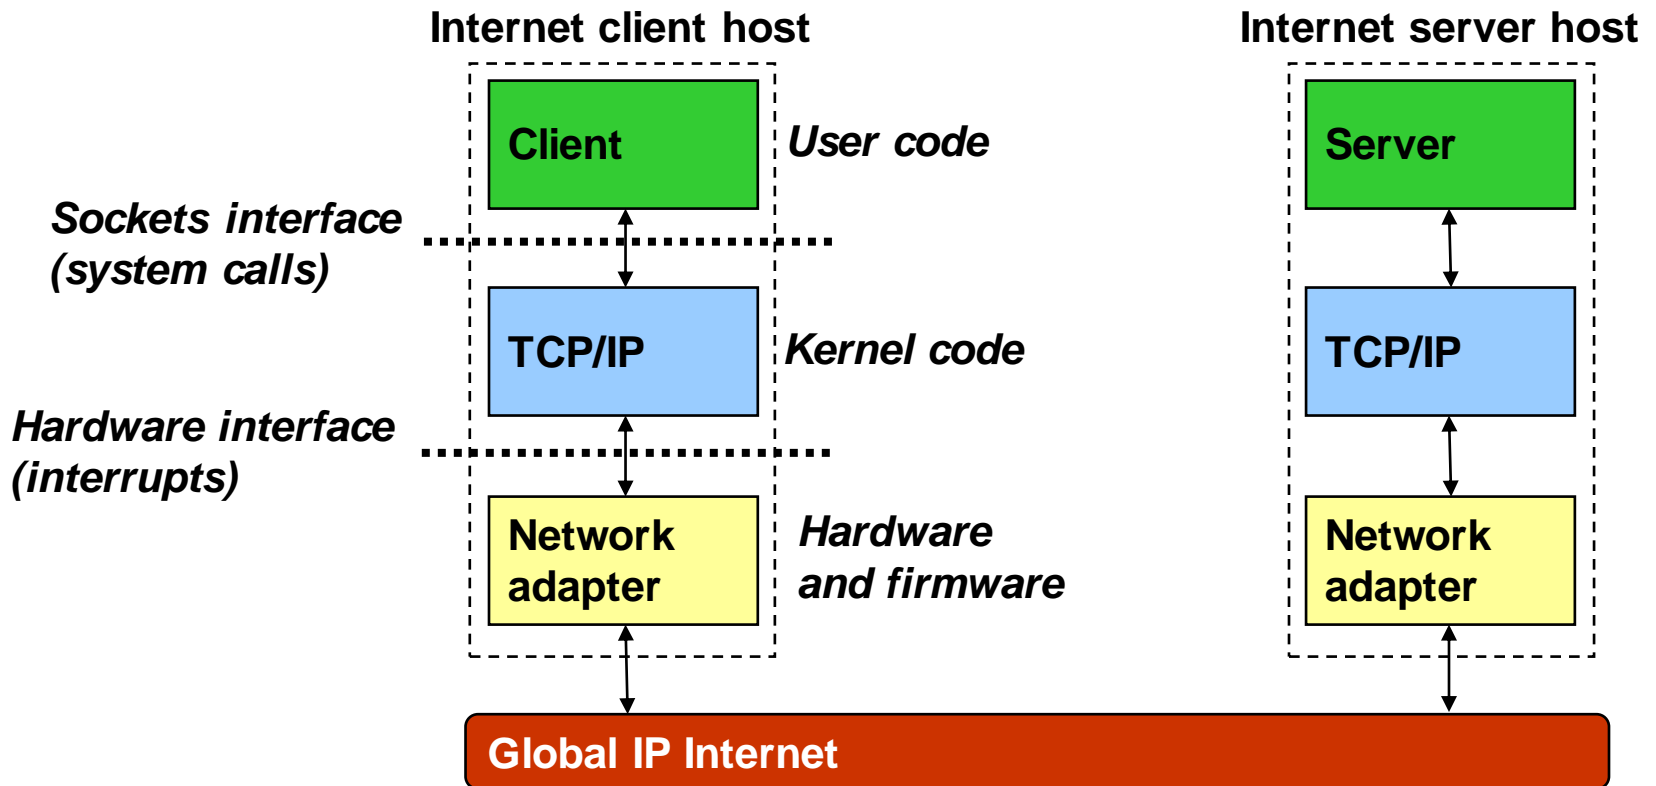


The Internet

- Most famous example of an internet
- Based on the TCP/IP protocol family
 - IP (Internet Protocol):
 - Provides *basic naming scheme* and *unreliable delivery capability* of packets (datagrams) from *host-to-host*
 - UDP (Unreliable Datagram Protocol)
 - Uses IP to provide *unreliable* datagram delivery from *process-to-process*
 - TCP (Transmission Control Protocol)
 - Uses IP to provide *reliable* byte streams from *process-to-process* over *connections*
- Accessed via a mix of Unix file I/O and functions from the *sockets interface*

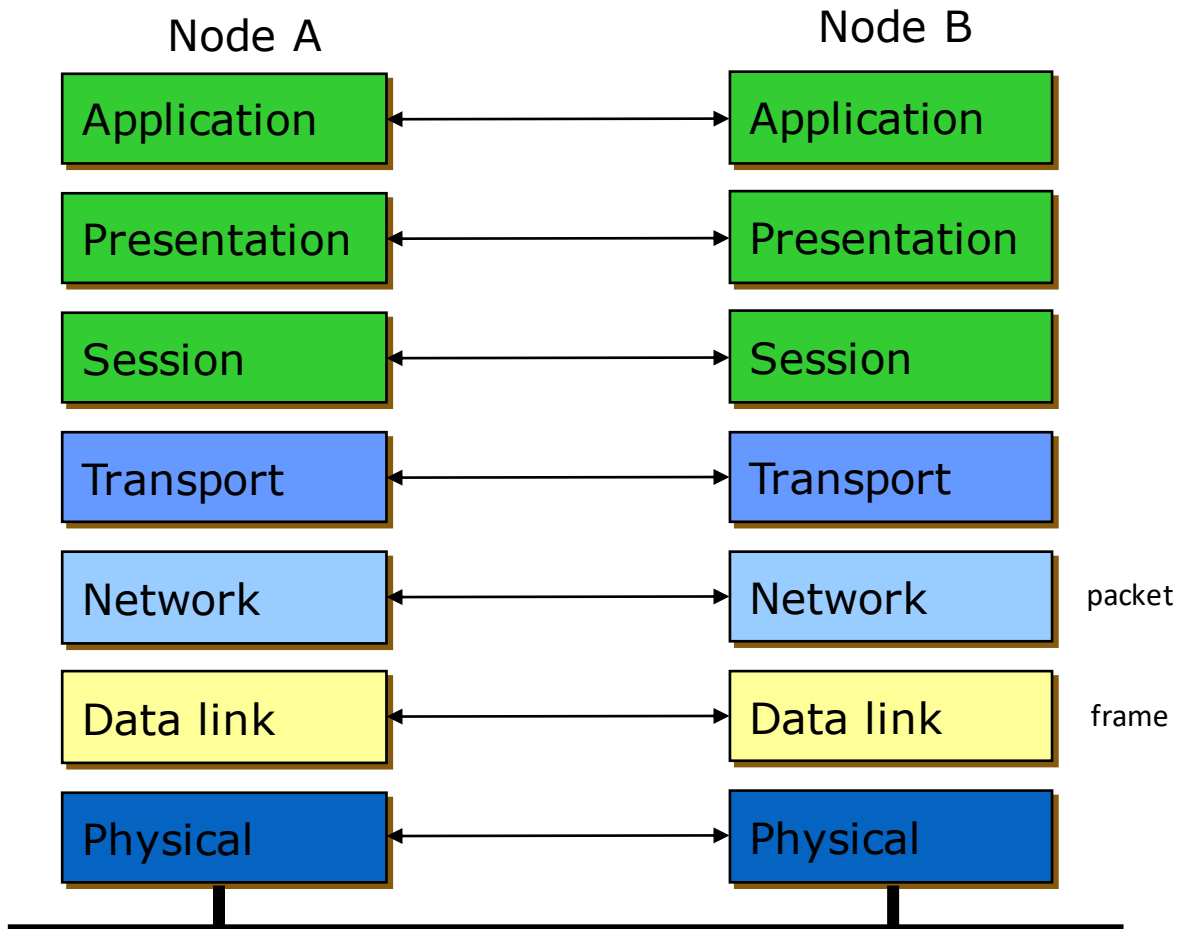
Organization of an Internet Application

- Hardware + Software



OSI reference model

- OSI (open system interconnection model)



A Programmer's View of the Internet

(1) Hosts are mapped to a set of 32-bit *IP addresses*

- 115.145.129.40

(2) The set of IP addresses is mapped to a set of identifiers called Internet *domain names*

- 115.145.129.40 is mapped to `www.skku.edu`

(3) A process on one Internet host can communicate with a process on another Internet host over a *connection*

Aside: IPv4 and IPv6

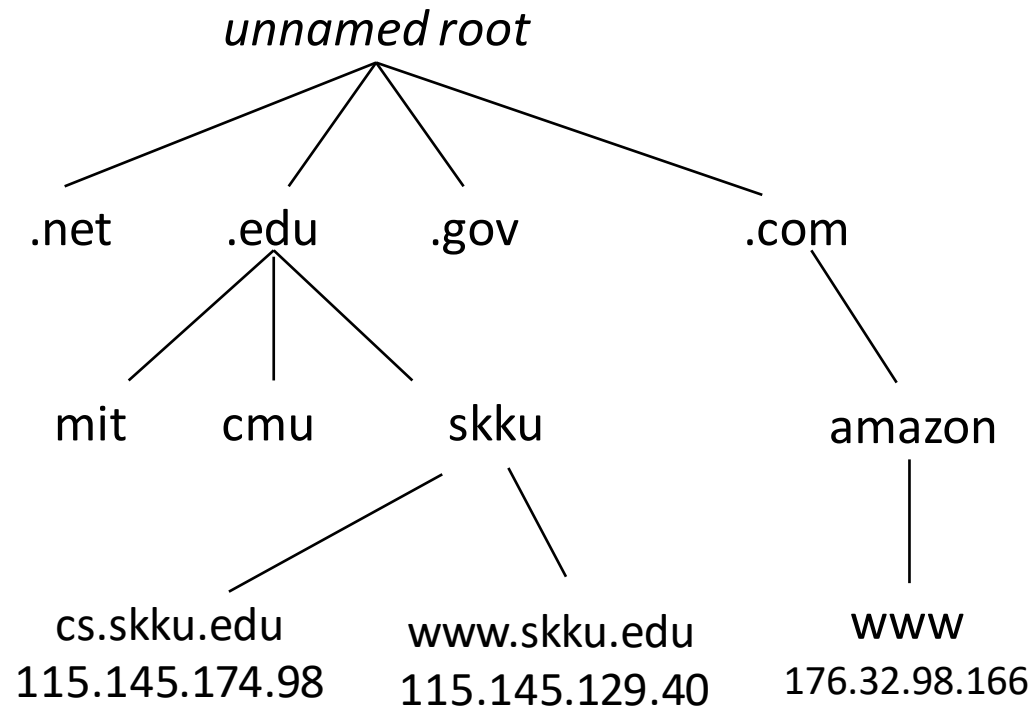
- The original Internet Protocol, with its 32-bit addresses, is known as *Internet Protocol Version 4 (IPv4)*
- 1996: Internet Engineering Task Force (IETF) introduced *Internet Protocol Version 6 (IPv6)* with 128-bit addresses
 - Intended as the successor to IPv4
 - As of 2015, vast majority of Internet traffic still carried by IPv4

(1) IP Addresses

- 32-bit IP addresses are stored in an *IP address struct*
 - IP addresses are always stored in memory in *network byte order* (big-endian byte order)
 - True in general for any integer transferred in a packet header from one machine to another.
 - E.g., the port number used to identify an Internet connection.
- By convention, each byte in a 32-bit IP address is represented by its decimal value and separated by a period
 - IP address: 0x8002C2F2 = 128.2.194.242

(2) Internet Domain Names

- More human-friendly domain names



Domain Naming System (DNS)

- The Internet maintains a mapping between IP addresses and domain names in a huge worldwide distributed database called *DNS*
- Conceptually, programmers can view the DNS database as a collection of millions of *host entries*.
 - Each host entry defines the mapping between a set of domain names and IP addresses.
 - In a mathematical sense, a host entry is an equivalence class of domain names and IP addresses.

(3) Internet Connections

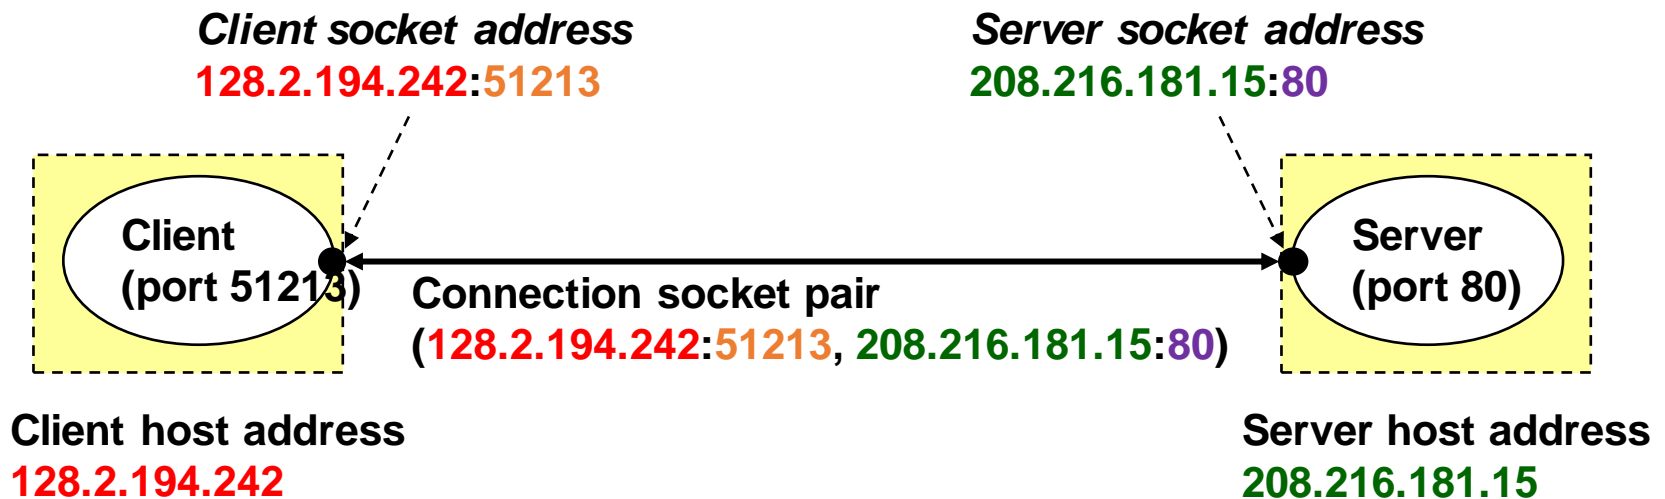
- Clients and servers communicate by sending streams of bytes over *connections*. Each connection is:
 - *Point-to-point*: connects a pair of processes.
 - *Full-duplex*: data can flow in both directions at the same time,
 - *Reliable*: stream of bytes sent by the source is eventually received by the destination in the same order it was sent.
- A *socket* is an endpoint of a connection
 - *Socket address* is an **IPAddress:port** pair
- A *port* is a 16-bit integer that identifies a process:
 - *Ephemeral port*: Assigned automatically by client kernel when client makes a connection request.
 - *Well-known port*: Associated with some *service* provided by a server (e.g., port 80 is associated with Web servers)

Well-known Ports and Service Names

- Popular services have permanently assigned *well-known ports* and corresponding *well-known service names*:
 - echo server: 7/echo
 - ssh servers: 22/ssh
 - email server: 25/smtp
 - Web servers: 80/http
- Mappings between well-known ports and service names is contained in the file `/etc/services` on each Linux machine.

Internet Connections

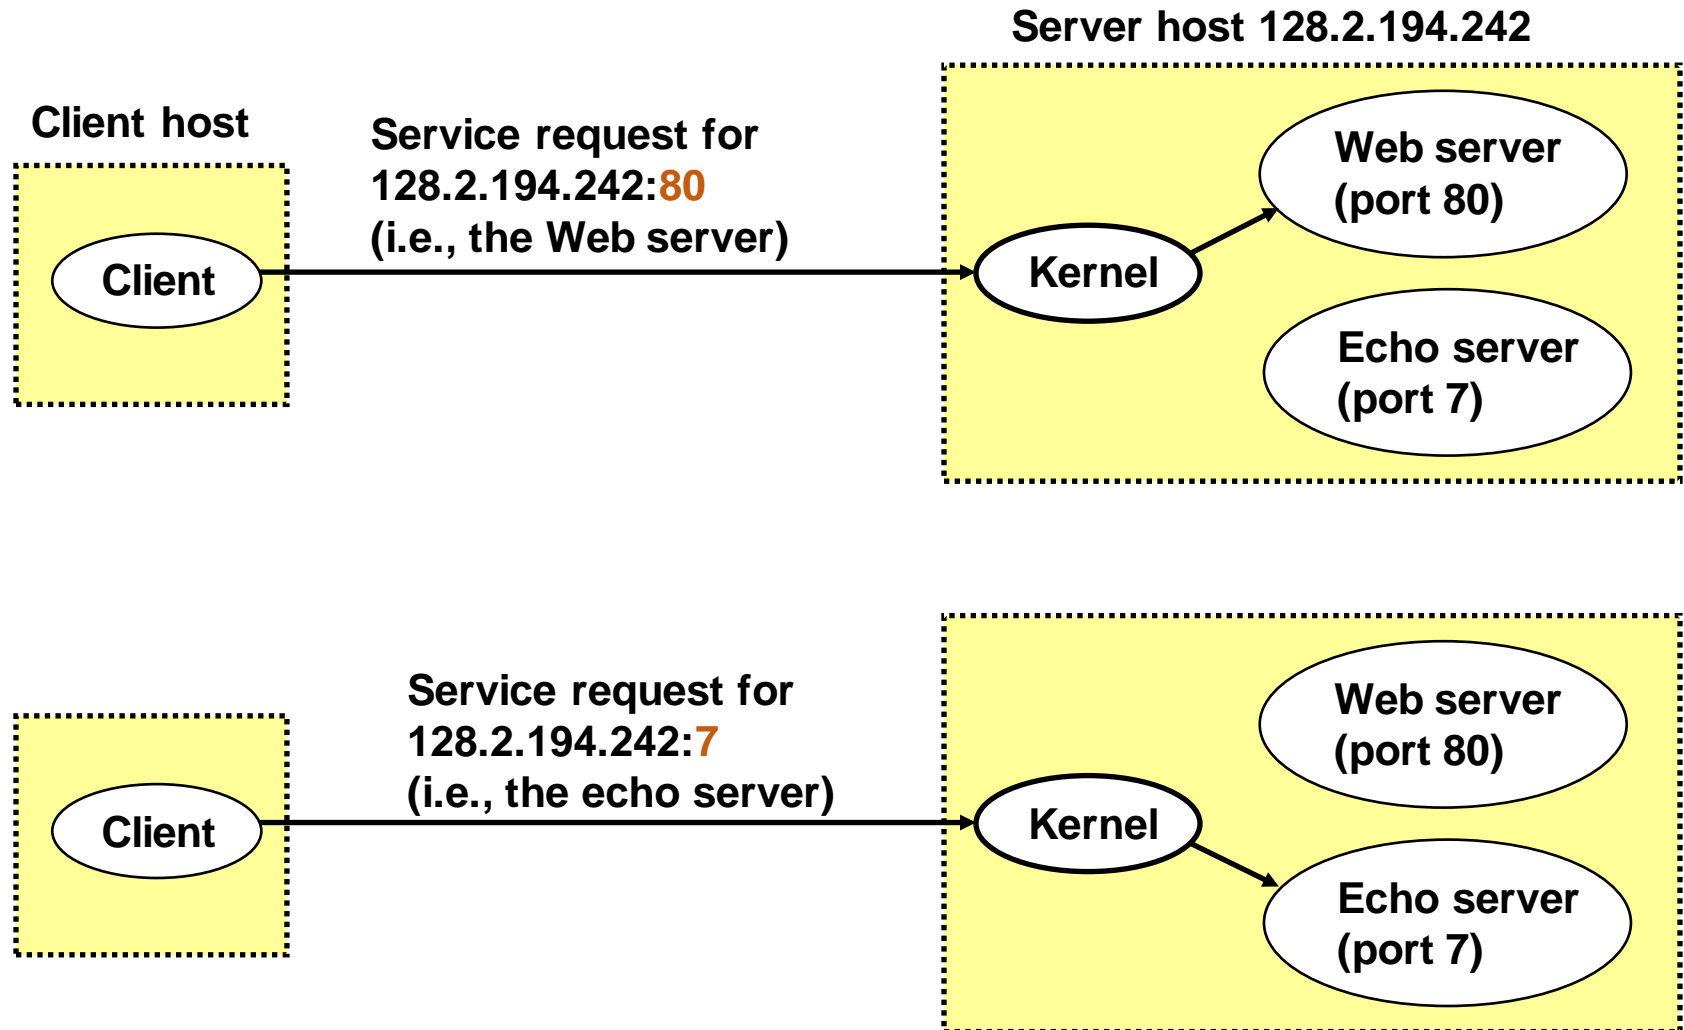
- A connection is uniquely identified by the socket addresses of its endpoints (*socket pair*)
 - (cliaddr:cliport, servaddr:servport)



Note: *51213* is an ephemeral port allocated by the kernel

Note: *80* is a well-known port associated with Web servers

Using Ports



Sockets Interface

- Set of **system-level functions** used in conjunction with Unix I/O to build network applications.
- Created in the early 80's as part of the original Berkeley distribution of Unix that contained an early version of the Internet protocols.
- Available on all modern systems
 - Unix variants, Windows, OS X, IOS, Android, ARM

Sockets


- What is a socket?
 - To the kernel, a socket is an **endpoint** of communication (connection)
 - To an application, a socket is a **file descriptor (an open file w/ a corresponding descriptor)** that lets the application read/write from/to the network
 - **Remember:** All Unix I/O devices, including networks, are modeled as files
- Clients and servers communicate with each other by reading from and writing to socket descriptors



- The main distinction between regular file I/O and socket I/O is how the application “opens” the socket descriptors


Sockets Interface

2. Start client *Client*



open_clientfd


1. Start server *Server*



open_listenfd

Connection request


Client/
Server
Session




3. Exchange data

Await connection request from next client

4. Disconnect client



EOF



5. Drop client

Discussion : Explain how system calls are used to implement this echo server

Echo Server (1)

```
#include <stdio.h>
#include <stdlib.h>
#include <sys/types.h>
#include <sys/socket.h>
#include <netdb.h>
#define MAXLINE 80

int main (int argc, char *argv[]) {
    int n, listenfd, connfd, caddrlen;
    struct hostent *h;
    struct sockaddr_in saddr, caddr;
    char buf[MAXLINE];
    int port = atoi(argv[1]);

    if ((listenfd = socket(AF_INET, SOCK_STREAM, 0)) < 0) {
        printf("socket() failed.\n");
        exit(1);
    }

    bzero((char *)&saddr, sizeof(saddr));
    saddr.sin_family = AF_INET;
    saddr.sin_addr.s_addr = htonl(INADDR_ANY);
    saddr.sin_port = htons(port);
```

Echo Server (2)

```
if (bind(listenfd, (struct sockaddr *)&saddr,  
          sizeof(saddr)) < 0) {  
    printf("bind() failed.\n");  
    exit(2);  
}  
  
if (listen(listenfd, 5) < 0) {  
    printf("listen() failed.\n");  
    exit(3);  
}  
  
while (1) {  
    caddrlen = sizeof(caddr);  
    if ((connfd = accept(listenfd, (struct sockaddr *)&caddr,  
                        &caddrlen)) < 0) {  
        printf ("accept() failed.\n");  
        continue;  
    }  
}
```

Echo Server (3)

```
// echo
while ((n = read(connfd, buf, MAXLINE)) > 0) {
    printf ("got %d bytes from client.\n", n);
    write(connfd, buf, n);
}

printf("connection terminated.\n");
close(connfd);
}
}
```